



# Computational Science on Many-Core Architectures

360.252

**Karl Rupp**



Institute for Microelectronics  
Vienna University of Technology  
<https://www.iue.tuwien.ac.at/>



Zoom Channel 95028746244  
Wednesday, November 11, 2020

# Agenda for Today

---


Exercise 3 Recap

Kernel Fusion - Conjugate Gradients

Kernel Fusion - Multiple Dot Products

Exercise 4

## Exercise 3 Recap



<https://xkcd.com/303/>

# Exercise 3 Recap

---

## Feedback Time

- How was your experience?

## Exercise 3 Recap

---

### Feedback Time

- How was your experience?
- Links to points for course will be sent out today or tomorrow

# Exercise 3 Recap

## Feedback Time

- How was your experience?
- Links to points for course will be sent out today or tomorrow

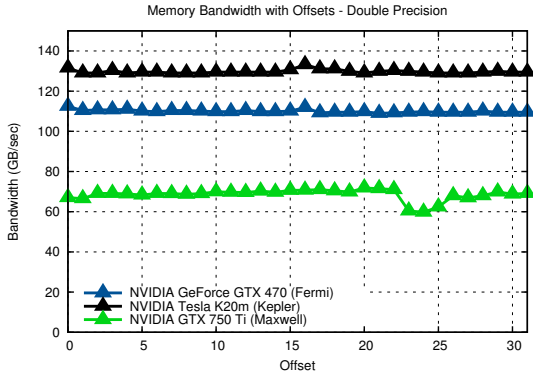
## Mind the Temporaries!

```
double dot_product(int N, ...) {  
    double result, *cuda_result;  
    cudaMalloc(&cuda_result, sizeof(double));  
    cuda_dot_product<<<...>>>(...);  
    ...  
}  
  
void cg() {  
    ...  
    while (1) {  
        ...  
        dot_product(...);  
    }  
}
```

# Exercise 3 Recap

## Offset Memory Access


```
__global__  
void work(double *x, double *y, double *z, int N, int k)  
{  
    int thread_id = blockIdx.x*blockDim.x + threadIdx.x;  
    for (size_t i=thread_id; i<N; i += blockDim.x * gridDim.x)  
        z[i+k] = x[i+k] + y[i+k];  
}
```



# Exercise 3 Recap

## Strided Memory Access

```
__global__  
void work(double *x, double *y, double *z, int N, int k)  
{  
    int thread_id = blockIdx.x*blockDim.x + threadIdx.x;  
    for (size_t i=thread_id; i<N; i += blockDim.x * gridDim.x)  
        z[i*k] = x[i*k] + y[i*k];  
}
```





# Exercise 3 Recap

## Strided Memory Access

- Array of structs problematic

```
typedef struct particle
{
    double pos_x; double pos_y; double pos_z;
    double vel_x; double vel_y; double vel_z;
    double mass;
} Particle;

__global__
void increase_mass(Particle *particles, int N)
{
    int thread_id = blockIdx.x*blockDim.x + threadIdx.x;
    for (int i=thread_id; i<N; i += blockDim.x * gridDim.x)
        particles[i].mass *= 2.0;
}
```

# Exercise 3 Recap

## Strided Memory Access

- Workaround: Structure of Arrays

```
typedef struct particles
{
    double *pos_x; double *pos_y; double *pos_z;
    double *vel_x; double *vel_y; double *vel_z;
    double *mass;
} Particle;

__global__
void increase_mass(Particle *particles, int N)
{
    int thread_id = blockIdx.x*blockDim.x + threadIdx.x;
    for (int i=thread_id; i<N; i += blockDim.x * gridDim.x)
        particles.mass[i] *= 2.0;
}
```

# Conjugate Gradients

## Pseudocode

Choose  $x_0$

$$p_0 = r_0 = b - Ax_0$$

For  $i = 0$  until convergence

1. Compute and store  $Ap_i$
2. Compute  $\langle p_i, Ap_i \rangle$
3.  $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
4.  $x_{i+1} = x_i + \alpha_i p_i$
5.  $r_{i+1} = r_i - \alpha_i Ap_i$
6. Compute  $\langle r_{i+1}, r_{i+1} \rangle$
7.  $\beta_i = \langle r_{i+1}, r_{i+1} \rangle / \langle r_i, r_i \rangle$
8.  $p_{i+1} = r_{i+1} + \beta_i p_i$

EndFor

## BLAS-based Implementation

-

SpMV, AXPY

For  $i = 0$  until convergence

1. SpMV
2. DOT
3. -
4. AXPY
5. AXPY
6. DOT
7. -
8. AXPY

EndFor

# Conjugate Gradients

## Pseudocode

Choose  $x_0$

$$p_0 = r_0 = b - Ax_0$$

For  $i = 0$  until convergence

1. Compute and store  $Ap_i$
2. Compute  $\langle p_i, Ap_i \rangle$
3.  $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
4.  $x_{i+1} = x_i + \alpha_i p_i$
5.  $r_{i+1} = r_i - \alpha_i Ap_i$
6. Compute  $\langle r_{i+1}, r_{i+1} \rangle$
7.  $\beta_i = \langle r_{i+1}, r_{i+1} \rangle / \langle r_i, r_i \rangle$
8.  $p_{i+1} = r_{i+1} + \beta_i p_i$

EndFor

## BLAS-based Implementation

-

SpMV, AXPY

For  $i = 0$  until convergence

1. SpMV
2. DOT  $\leftarrow$  Global sync!
3. -
4. AXPY
5. AXPY
6. DOT  $\leftarrow$  Global sync!
7. -
8. AXPY

EndFor

# Conjugate Gradients

## Pseudocode

Choose  $x_0$

$$p_0 = r_0 = b - Ax_0$$

For  $i = 0$  until convergence

1. Compute and store  $Ap_i$
2. Compute  $\langle p_i, Ap_i \rangle$
3.  $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
4.  $x_{i+1} = x_i + \alpha_i p_i$
5.  $r_{i+1} = r_i - \alpha_i Ap_i$
6. Compute  $\langle r_{i+1}, r_{i+1} \rangle$
7.  $\beta_i = \langle r_{i+1}, r_{i+1} \rangle / \langle r_i, r_i \rangle$
8.  $p_{i+1} = r_{i+1} + \beta_i p_i$

EndFor

## BLAS-based Implementation

-


SpMV, AXPY

For  $i = 0$  until convergence

1. SpMV  $\leftarrow$  No caching of  $Ap_i$
2. DOT  $\leftarrow$  Global sync!
3. -
4. AXPY
5. AXPY  $\leftarrow$  No caching of  $r_{i+1}$
6. DOT  $\leftarrow$  Global sync!
7. -
8. AXPY

EndFor

# Conjugate Gradients



- (2D Finite Difference Discretization)

# Conjugate Gradient

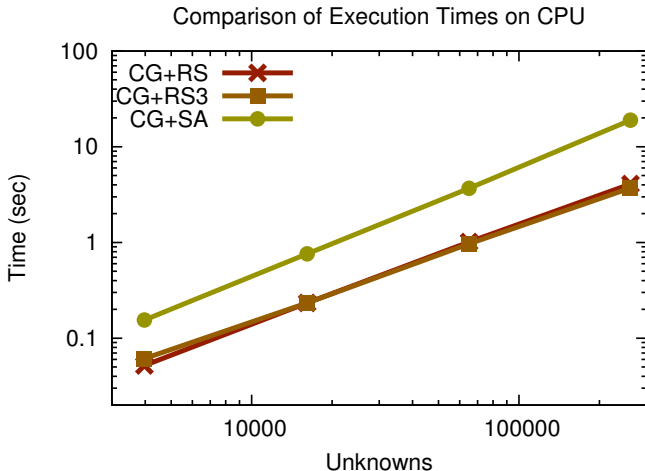
## Implications

- Kernel launches expensive
- Delicate balance for preconditioners

# Conjugate Gradient

## Implications

- Kernel launches expensive
- Delicate balance for preconditioners






# Conjugate Gradient

## Implications

- Kernel launches expensive
- Delicate balance for preconditioners




# Conjugate Gradient


## Implications

- Kernel launches expensive
- Delicate balance for preconditioners

Comparison of Execution Times on CPU



Comparison of Execution Times on GPU




# Conjugate Gradient


## Implications

- Kernel launches expensive
- Delicate balance for preconditioners

Comparison of Execution Times on CPU



Comparison of Execution Times on GPU




# Conjugate Gradient Optimizations

## Optimization 1

- Get best performance out of SpMV
- Compare different sparse matrix types


Cf.: N. Bell: Implementing sparse matrix-vector multiplication on throughput-oriented processors. *Proc. SC '09*

# Conjugate Gradients




- (2D Finite Difference Discretization)

# Conjugate Gradients



- (2D Finite Difference Discretization)

# Conjugate Gradients



- (2D Finite Difference Discretization)

# Conjugate Gradient Optimizations


---

## Optimization 2

- Optimize kernel parameters for each operation




# Conjugate Gradients



- (2D Finite Difference Discretization)

# Conjugate Gradients



- (2D Finite Difference Discretization)

# Conjugate Gradient Optimizations

## Optimization 3: Rearrange the algorithm

- Remove unnecessary reads
- Remove unnecessary synchronizations
- Use custom kernels instead of standard BLAS

# Conjugate Gradients

## Standard CG

Choose  $x_0$

$$p_0 = r_0 = b - Ax_0$$

For  $i = 0$  until convergence

1. Compute and store  $Ap_i$
2. Compute  $\langle p_i, Ap_i \rangle$
3.  $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
4.  $x_{i+1} = x_i + \alpha_i p_i$
5.  $r_{i+1} = r_i - \alpha_i Ap_i$
6. Compute  $\langle r_{i+1}, r_{i+1} \rangle$
7.  $\beta_i = \langle r_{i+1}, r_{i+1} \rangle / \langle r_i, r_i \rangle$
8.  $p_{i+1} = r_{i+1} + \beta_i p_i$

EndFor

# Conjugate Gradients

## Standard CG

Choose  $x_0$

$$p_0 = r_0 = b - Ax_0$$

For  $i = 0$  until convergence

1. Compute and store  $Ap_i$
2. Compute  $\langle p_i, Ap_i \rangle$
3.  $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
4.  $x_{i+1} = x_i + \alpha_i p_i$
5.  $r_{i+1} = r_i - \alpha_i Ap_i$
6. Compute  $\langle r_{i+1}, r_{i+1} \rangle$
7.  $\beta_i = \langle r_{i+1}, r_{i+1} \rangle / \langle r_i, r_i \rangle$
8.  $p_{i+1} = r_{i+1} + \beta_i p_i$

EndFor

## Pipelined CG

Choose  $x_0$

$$p_0 = r_0 = b - Ax_0$$

For  $i = 1$  until convergence

1.  $i = 1$ : Compute  $\alpha_0, \beta_0, Ap_0$
2.  $x_i = x_{i-1} + \alpha_{i-1} p_{i-1}$
3.  $r_i = r_{i-1} - \alpha_{i-1} Ap_{i-1}$
4.  $p_i = r_i + \beta_{i-1} p_{i-1}$
5. Compute and store  $Ap_i$
6. Compute  $\langle Ap_i, Ap_i \rangle, \langle p_i, Ap_i \rangle, \langle r_i, r_i \rangle$
7.  $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
8.  $\beta_i = (\alpha_i^2 \langle Ap_i, Ap_i \rangle - \langle r_i, r_i \rangle) / \langle r_i, r_i \rangle$

EndFor

# Conjugate Gradients

## Standard CG

Choose  $x_0$

$$p_0 = r_0 = b - Ax_0$$

For  $i = 0$  until convergence

1. Compute and store  $Ap_i$
2. Compute  $\langle p_i, Ap_i \rangle$
3.  $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
4.  $x_{i+1} = x_i + \alpha_i p_i$
5.  $r_{i+1} = r_i - \alpha_i Ap_i$
6. Compute  $\langle r_{i+1}, r_{i+1} \rangle$
7.  $\beta_i = \langle r_{i+1}, r_{i+1} \rangle / \langle r_i, r_i \rangle$
8.  $p_{i+1} = r_{i+1} + \beta_i p_i$

EndFor

## Pipelined CG

Choose  $x_0$

$$p_0 = r_0 = b - Ax_0$$

For  $i = 1$  until convergence

1.  $i = 1$ : Compute  $\alpha_0, \beta_0, Ap_0$
2.  $x_i = x_{i-1} + \alpha_{i-1} p_{i-1}$
3.  $r_i = r_{i-1} - \alpha_{i-1} Ap_{i-1}$
4.  $p_i = r_i + \beta_{i-1} p_{i-1}$
5. Compute and store  $Ap_i$
6. Compute  $\langle Ap_i, Ap_i \rangle, \langle p_i, Ap_i \rangle, \langle r_i, r_i \rangle$
7.  $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
8.  $\beta_i = (\alpha_i^2 \langle Ap_i, Ap_i \rangle - \langle r_i, r_i \rangle) / \langle r_i, r_i \rangle$

EndFor

# Conjugate Gradients

## Standard CG

Choose  $x_0$

$$p_0 = r_0 = b - Ax_0$$

For  $i = 0$  until convergence

1. Compute and store  $Ap_i$
2. Compute  $\langle p_i, Ap_i \rangle$
3.  $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
4.  $x_{i+1} = x_i + \alpha_i p_i$
5.  $r_{i+1} = r_i - \alpha_i Ap_i$
6. Compute  $\langle r_{i+1}, r_{i+1} \rangle$
7.  $\beta_i = \langle r_{i+1}, r_{i+1} \rangle / \langle r_i, r_i \rangle$
8.  $p_{i+1} = r_{i+1} + \beta_i p_i$

EndFor

## Pipelined CG

Choose  $x_0$


$$p_0 = r_0 = b - Ax_0$$

For  $i = 1$  until convergence

1.  $i = 1$ : Compute  $\alpha_0, \beta_0, Ap_0$
2.  $x_i = x_{i-1} + \alpha_{i-1} p_{i-1}$
3.  $r_i = r_{i-1} - \alpha_{i-1} Ap_{i-1}$
4.  $p_i = r_i + \beta_{i-1} p_{i-1}$
5. Compute and store  $Ap_i$
6. Compute  $\langle Ap_i, Ap_i \rangle, \langle p_i, Ap_i \rangle, \langle r_i, r_i \rangle$
7.  $\alpha_i = \langle r_i, r_i \rangle / \langle p_i, Ap_i \rangle$
8.  $\beta_i = (\alpha_i^2 \langle Ap_i, Ap_i \rangle - \langle r_i, r_i \rangle) / \langle r_i, r_i \rangle$

EndFor


# Conjugate Gradients



- (2D Finite Difference Discretization)




# Conjugate Gradients




- (2D Finite Difference Discretization)

# Conjugate Gradients



- (2D Finite Difference Discretization)

# Conjugate Gradients



- (2D Finite Difference Discretization)

# Kernel Fusion - Multiple Dot Products

## Gram-Schmidt method

- Given orthonormal basis  $\{v_1, v_2, \dots, v_N\}$ , augment by  $w$
- $w \leftarrow w - \langle w, v_i \rangle v_i$
- $w \leftarrow w / \|w\|$
- Add  $w$  to basis

## Multiple inner products $\langle w, v_i \rangle$


- Performance critical (global reductions)
- Reuse of  $w$  desirable

# Implementation 1: dot

Goal: Compute  $\alpha_i = \langle w, v_i \rangle$  for  $i = 1, \dots, N$  efficiently

## Method 1: Iterated application of *dot()*

- Decompose into  $N$  separate inner products
- $N$  calls to BLAS level 1 routine *ddot*
- Reductions: One per work group, one over results of work groups



## Implementation 2: gemv

Goal: Compute  $\alpha_i = \langle w, v_i \rangle$  for  $i = 1, \dots, N$  efficiently

Method 2: Pack vectors into matrix, use *gemv*

- Set  $\mathbf{A} = \begin{pmatrix} v_1^T \\ \vdots \\ v_N^T \end{pmatrix} \in \mathbb{R}^{N \times M}, N \ll M$
- Compute  $\alpha = \mathbf{A}\mathbf{x}$
- One BLAS level 2 *dgemv* call

## Implementation 3: mdot

### Method 3: Custom routine *mdot*

- Process  $\alpha_i = \langle w, v_i \rangle$  in batches
- Batch sizes 1, 2, 3, 4, 8
- Load entries of  $w$  only once per batch
- Fit remaining inner products into largest batch possible
- Batch size 8: Only 12.5% overhead vs. arbitrary batch sizes

# Implementation 3: mdot

## Method 3: Custom routine *mdot*

- Process  $\alpha_i = \langle w, v_i \rangle$  in batches
- Batch sizes 1, 2, 3, 4, 8
- Load entries of  $w$  only once per batch
- Fit remaining inner products into largest batch possible
- Batch size 8: Only 12.5% overhead vs. arbitrary batch sizes

## Code sketch (Batch size 4)

```
for (size_t i=thread_id; i<M; i += threads_per_group)
{
    double val_w = w[i];
    alpha_1 += val_w * v1[i];
    alpha_2 += val_w * v2[i];
    alpha_3 += val_w * v3[i];
    alpha_4 += val_w * v4[i];
}
```



# Benchmarks

---

## Part 3: Benchmarks

# Exercises

## Environment

- <https://gtx1080.360252.org/2020/ex4/>
- (Might receive visual updates and additional hints over the next days)
- Due: Tuesday, November 17, 2020 at 23:59pm

## Hints and Suggestions

- Consider version control for locally developed code
- Please let me know of any bugs or issues
- Example codes and code skeletons provided at URL above