



Computational Science on Many-Core Architectures

360.252

Karl Rupp



Institute for Microelectronics
Vienna University of Technology
<http://www.iue.tuwien.ac.at>



Zoom Channel 95028746244
Wednesday, November 18, 2020

Agenda for Today

Exercise 4 Recap

Finite Differences

Recap: Prefix Sums

Sequential to Parallel: ILU Example

Exercise 5

Exercise 4 Recap

Feedback Time

- How was your experience?

Exercise 4 Recap

Feedback Time

- How was your experience?
- You should have received links to points for course last week

Exercise 4 Recap

Feedback Time

- How was your experience?
- You should have received links to points for course last week

Vector Iterations

```
__kernel__ void option1(double *x, int N, ...) {  
    for (int row = blockDim.x * blockIdx.x + threadIdx.x;  
         row < N;  
         row += gridDim.x * blockDim.x) { ... }  
}  
option1<<<256, 256>>>(x, N, ...); // arbitrary grid/block dim
```

```
__kernel__ void option2(double *x, int N, ...) {  
    if (blockDim.x*blockIdx.x + threadIdx.x < N) { ... }  
}  
option2<<<N/blocksize + 1, blocksize>>>(x, N, ...); // beware!
```

Discrete Poisson Equation

- Rectangular grid
- Assume homogeneous grid spacing h
- Discretize each coordinate direction separately

$$\begin{aligned}\Delta u(x, y) &= u_{xx}(x, y) + u_{yy}(x, y) \\ &\approx \frac{u(x-h, y) - 2u(x, y) + u(x+h, y)}{h^2} + \frac{u(x, y-h) - 2u(x, y) + u(x, y+h)}{h^2} \\ &= \frac{u(x-h, y) + u(x+h, y) - 4u(x, y) + u(x, y-h) + u(x, y+h)}{h^2} \\ &=: \Delta_h u(x, y)\end{aligned}$$

Finite Differences

System Matrix Structure

- Interior nodes have 5 nonzero entries
- Boundary nodes have 1 to 4 nonzero entries
- Problem: CSR format requires offset index for each row

	2		7	4
		1	9	
3				
		6		5
	8			

Matrix

2	7	4	1	9	3	6	5	8
1	3	4	2	3	0	2	4	1
0	3	5	6	8	9			

CSR

Finite Differences

System Matrix Structure

- Interior nodes have 5 nonzero entries
- Boundary nodes have 1 to 4 nonzero entries
- Problem: CSR format requires offset index for each row

	2		7	4
		1	9	
3				
		6		5
	8			

Matrix

2	7	4	1	9	3	6	5	8
1	3	4	2	3	0	2	4	1
0	3	5	6	8	9			

CSR

Generic GPU matrix assembly skeleton

1. Count the nonzero entries for each row
2. Deduce offsets for CSR
3. Populate column and values arrays

Finite Differences

Step 1: Count Nonzeros for an $N \times M$ grid

```
__kernel__ void count_nnz(int *row_offsets, int N, ...) {  
    for (int row = blockDim.x * blockIdx.x + threadIdx.x;  
         row < N*M;  
         row += gridDim.x * blockDim.x)  
    {  
        int nnz_for_this_node = 1;  
        int i = row / N;  
        int j = row % N;  
  
        if (i > 0) nnz_for_this_node += 1;  
        if (j > 0) nnz_for_this_node += 1;  
        if (i < N-1) nnz_for_this_node += 1;  
        if (j < M-1) nnz_for_this_node += 1;  
  
        row_offsets[row] = nnz_for_this_node;  
    }  
}
```

Parallel Primitives

Prefix Sum


- Inclusive: Determine $y_i = \sum_{k=1}^i x_k$
- Exclusive: Determine $y_i = \sum_{k=1}^{i-1} x_k$, $y_1 = 0$

Example

- x: 4, 3, 6, 5, 4, 7, 4, 4, 4
- y: 4, 7, 13, 18, 22, 29, 33, 37, 41 (inclusive)
- y: 0, 4, 7, 13, 18, 22, 29, 33, 37 (exclusive)


Applications

- Sparse matrix setup
- Graph algorithms



Parallel Primitives

Prefix Sum Implementation



```
for(int stride = 1; stride < blockDim.x; stride *= 2)
{
    __syncthreads();
    shared_m[threadIdx.x] = my_value;
    __syncthreads();
    if (threadIdx.x >= stride)
        my_value += shared_m[threadIdx.x - stride];
}
__syncthreads();
shared_m[threadIdx.x] = my_value;
```

Finite Differences

Step 3: Assembly for an $N \times M$ grid

```
__kernel__ void assembleA(int *row_offsets, int N, ...) {  
    for (int row = blockDim.x * blockIdx.x + threadIdx.x;  
         row < N*M;  
         row += gridDim.x * blockDim.x) {  
        int i = row / N;  
        int j = row % N;  
        int this_row_offset = row_offsets[row];  
  
        int index = i * N + j;  
        col_indices[this_row_offset] = index;  
        values[this_row_offset] = 4;  
        this_row_offset += 1;  
  
        if (i > 0) { /* similarly with correct index and value -1 */ }  
        if (j > 0) { /* similarly */ }  
        if (i < N-1) { /* similarly */ }  
        if (j < M-1) { /* similarly */ }  
    }  
}
```

Other Discretizations

Finite Volumes

- Iterations over vertices: Similar to finite differences
- Beware of advanced schemes (depends)


Other Discretizations

Finite Volumes

- Iterations over vertices: Similar to finite differences
- Beware of advanced schemes (depends)

Finite Elements

- Iteration over elements (cells)
- BUT: write matrix entries associated with vertices
- \Rightarrow Avoid race conditions
- More preparation required before we can address this



Sequential to Parallel: ILU Example

(in preparation for Finite Element assembly strategies)

Parallel ILU

ILU - Basic Idea

- Factor sparse matrix $A \approx \tilde{L}\tilde{U}$
- \tilde{L} and \tilde{U} sparse, triangular
- ILU0: Pattern of \tilde{L} , \tilde{U} equal to A
- ILUT: Keep k elements per row

Solver Cycle Phase

- Residual correction $\tilde{L}\tilde{U}x = z$
- Forward solve $\tilde{L}y = z$
- Backward solve $\tilde{U}x = y$
- Little parallelism in general

$$\begin{pmatrix} 5 & \times & \times & \times & & \times & \times & & \\ \times & 3 & \times & & & & & & \\ \times & \times & 4 & \times & & & & & \\ \times & & \times & 5 & \times & \times & & & \times \\ & & & \times & 5 & \times & & \times & \times \\ \times & & & \times & \times & 6 & \times & \times & \\ \times & & & & & \times & 3 & & \\ & & & & \times & \times & & 4 & \times \\ & & \times & \times & & & & \times & 4 \end{pmatrix}$$

Parallel ILU

ILU Level Scheduling

- Build dependency graph
- Substitute as many entries as possible simultaneously
- Trade-off: Each step vs. multiple steps in a single kernel

$$\begin{pmatrix} 5 & \times & \times & \times & & \times & \times & & \\ \times & 3 & \times & & & & & & \\ \times & \times & 4 & \times & & & & & \\ \times & & \times & 5 & \times & \times & & & \times \\ & & & \times & 5 & \times & & \times & \times \\ \times & & & \times & \times & 6 & \times & \times & \\ \times & & & & & \times & 3 & & \\ & & & & \times & \times & & 4 & \times \\ & & \times & \times & & & \times & \times & 4 \end{pmatrix}$$

Parallel ILU

ILU Level Scheduling

- Build dependency graph
- Substitute as many entries as possible simultaneously
- Trade-off: Each step vs. multiple steps in a single kernel

$$\begin{pmatrix} 5 & \times & \times & \times & & \times & \times & & \\ \times & 3 & \times & & & & & & \\ \times & \times & 4 & \times & & & & & \\ \times & & \times & 5 & \times & \times & & & \times \\ & & & \times & 5 & \times & & \times & \times \\ & & & \times & \times & 6 & \times & \times & \\ & & & & & \times & 3 & & \\ & & & & \times & \times & & 4 & \times \\ & & \times & \times & & & & \times & 4 \end{pmatrix}$$

Parallel ILU

ILU Level Scheduling



- Build dependency graph
- Substitute as many entries as possible simultaneously
- Trade-off: Each step vs. multiple steps in a single kernel

$$\begin{pmatrix} 5 & \times & \times & \times & & \times & \times & & \\ \boxed{\times} & 3 & \times & & & & & & \\ \times & \boxed{\times} & 4 & \times & & & & & \\ \times & & \boxed{\times} & 5 & \times & \times & & & \times \\ \times & & & \boxed{\times} & 5 & \times & & \times & \times \\ \times & & & \times & \times & 6 & \times & \times & \\ \times & & & & \times & \times & 3 & & \\ & & & \times & \times & & & 4 & \times \\ & & & \times & & & & \times & 4 \end{pmatrix}$$

Parallel ILU

ILU Interpretation on Structured Grids



- 2d finite-difference discretization
- Substitution whenever all neighbors with smaller index computed
- Works particularly well in 3d



Parallel ILU

ILU Interpretation on Structured Grids



- 2d finite-difference discretization
- Substitution whenever all neighbors with smaller index computed
- Works particularly well in 3d



Parallel ILU

ILU Interpretation on Structured Grids



- 2d finite-difference discretization
- Substitution whenever all neighbors with smaller index computed
- Works particularly well in 3d



Parallel ILU

ILU Interpretation on Structured Grids



- 2d finite-difference discretization
- Substitution whenever all neighbors with smaller index computed
- Works particularly well in 3d



Parallel ILU

ILU Interpretation on Structured Grids



- 2d finite-difference discretization
- Substitution whenever all neighbors with smaller index computed
- Works particularly well in 3d



Parallel ILU

ILU Interpretation on Structured Grids



- 2d finite-difference discretization
- Substitution whenever all neighbors with smaller index computed
- Works particularly well in 3d



Parallel ILU

ILU Interpretation on Structured Grids



- 2d finite-difference discretization
- Substitution whenever all neighbors with smaller index computed
- Works particularly well in 3d



Parallel ILU

ILU Interpretation on Structured Grids



- 2d finite-difference discretization
- Substitution whenever all neighbors with smaller index computed
- Works particularly well in 3d



Parallel ILU

ILU Interpretation on Structured Grids



- 2d finite-difference discretization
- Substitution whenever all neighbors with smaller index computed
- Works particularly well in 3d



Parallel ILU

ILU Interpretation on Structured Grids



- 2d finite-difference discretization
- Substitution whenever all neighbors with smaller index computed
- Works particularly well in 3d



Parallel ILU

ILU Interpretation on Structured Grids

- 2d finite-difference discretization
- Substitution whenever all neighbors with smaller index computed
- Works particularly well in 3d



Parallel ILU - Innovation

Sequential

```
for i=2..n
  for k=1..i-1, (i,k) in A
     $a_{ik} = a_{ik} / a_{kk}$ 
  for j=k+1..n, (i,j) in A
     $u_{ij} = a_{ij} - a_{ik}a_{kj}$ 
```

Parallel

```
for (sweep = 1, 2, ...)
  parallel for (i,j) in A
    if (i > j)
       $l_{ij} = (a_{ij} - \sum_{k=1}^{j-1} l_{ik}u_{kj}) / u_{jj}$ 
    else
       $u_{ij} = a_{ij} - \sum_{k=1}^{i-1} l_{ik}u_{kj}$ 
```

Fine-Grained Parallel ILU Setup

- Proposed by Chow and Patel (SISC, vol. 37(2)) for CPUs and MICs
- Massively parallel (one thread per row)

Parallel ILU - Innovation

Sequential

```
for i=2..n
  for k=1..i-1, (i,k) in A
     $a_{ik} = a_{ik} / a_{kk}$ 
  for j=k+1..n, (i,j) in A
     $u_{ij} = a_{ij} - a_{ik}a_{kj}$ 
```

Parallel

```
for (sweep = 1, 2, ...)
  parallel for (i,j) in A
    if (i > j)
       $l_{ij} = (a_{ij} - \sum_{k=1}^{j-1} l_{ik}u_{kj}) / u_{jj}$ 
    else
       $u_{ij} = a_{ij} - \sum_{k=1}^{i-1} l_{ik}u_{kj}$ 
```

Fine-Grained Parallel ILU Setup

- Proposed by Chow and Patel (SISC, vol. 37(2)) for CPUs and MICs
- Massively parallel (one thread per row)

Preconditioner Application

- Truncated Neumann series:

$$\mathbf{L}^{-1} \approx \sum_{k=0}^K (\mathbf{I} - \mathbf{L})^k, \quad \mathbf{U}^{-1} \approx \sum_{k=0}^K (\mathbf{I} - \mathbf{U})^k$$

- Exact triangular solves not necessary

Exercises

Environment

- <https://gtx1080.360252.org/2020/ex5/>
- (Might receive visual updates and additional hints over the next days)
- Due: Tuesday, November 24, 2020 at 23:59pm

Hints and Suggestions

- Consider version control for locally developed code
- Please let me know of any bugs or issues