



Computational Science on Many-Core Architectures

360.252

Karl Rupp



Institute for Microelectronics
Vienna University of Technology
<http://www.iue.tuwien.ac.at>



Zoom Channel 95028746244
Wednesday, December 2, 2020

Agenda for Today

Exercise 6 Recap

OpenCL

Exercise 7

Freakin' Fast Friday

Exercise 6 Recap

Feedback Time

- How was your experience?

Exercise 6 Recap

Feedback Time

- How was your experience?
- Points for Exercise 5 will be provided within 24 hours.

History of OpenCL

2008

- OpenCL working group formed at Khronos Group
- OpenCL specification 1.0 released

2010

- OpenCL 1.1 (multi-device, subbuffer)

2011

- OpenCL 1.2 (device partitioning)

2013

- OpenCL 2.0 (shared virtual memory, SPIR, etc.)

2020

- OpenCL 3.0 (back to the roots)



OpenCL

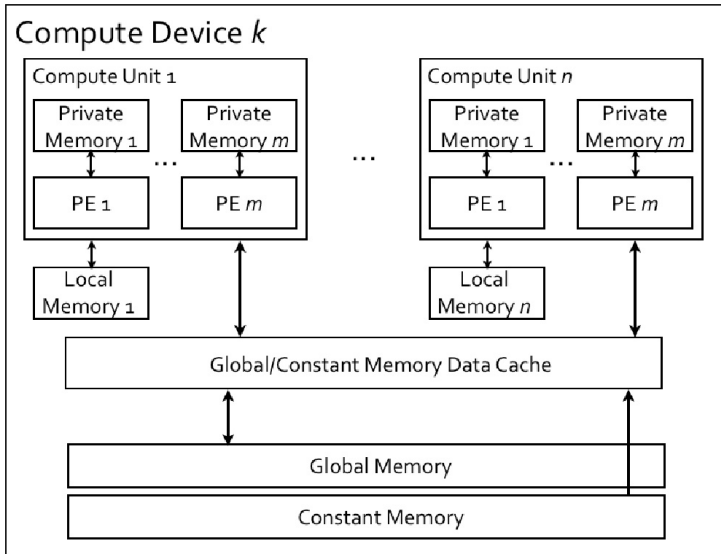
Similar to CUDA

- Kernel language is a subset of C
- Explicit memory management, host-device transfers
- Memory model: local, shared, global

Different from CUDA

- Support by many vendors
- No compiler-wrapper, only a shared library
- Kernel compilation usually at runtime


OpenCL Platform Model




OpenCL Platform Model

Platform 1

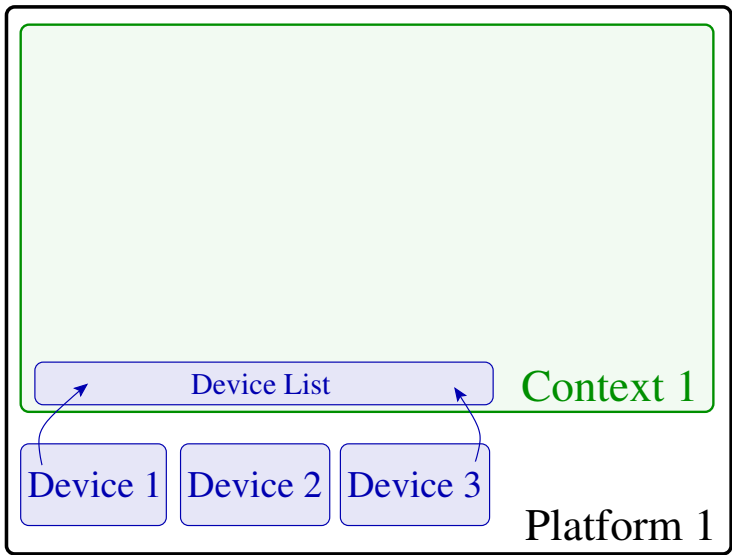
OpenCL Platform Model



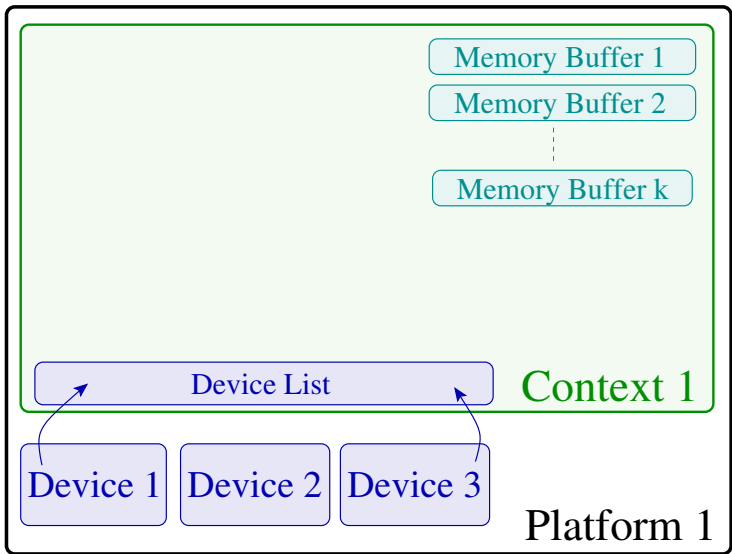
OpenCL Platform Model



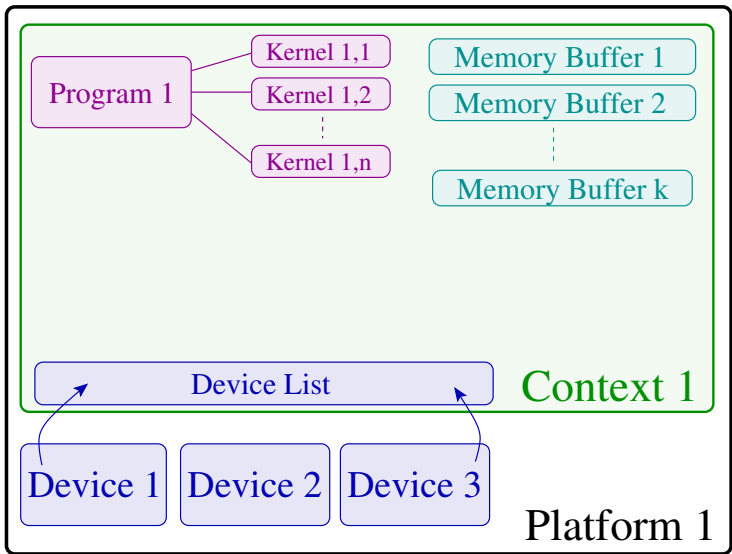
OpenCL Platform Model



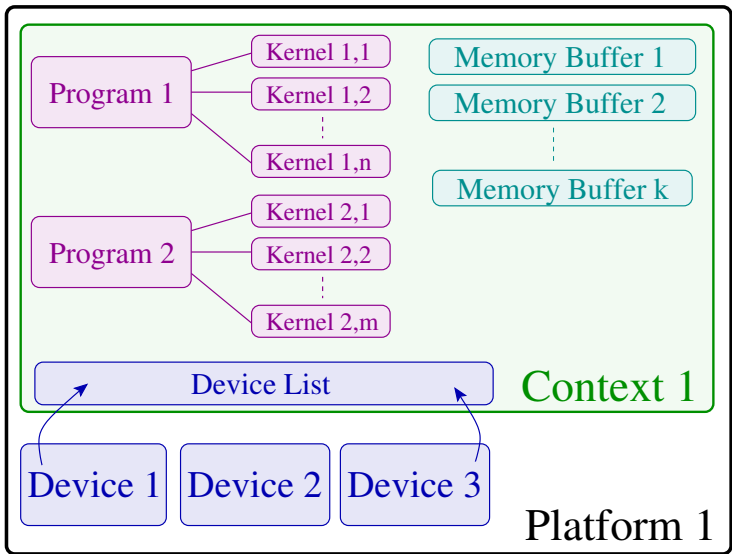
OpenCL Platform Model



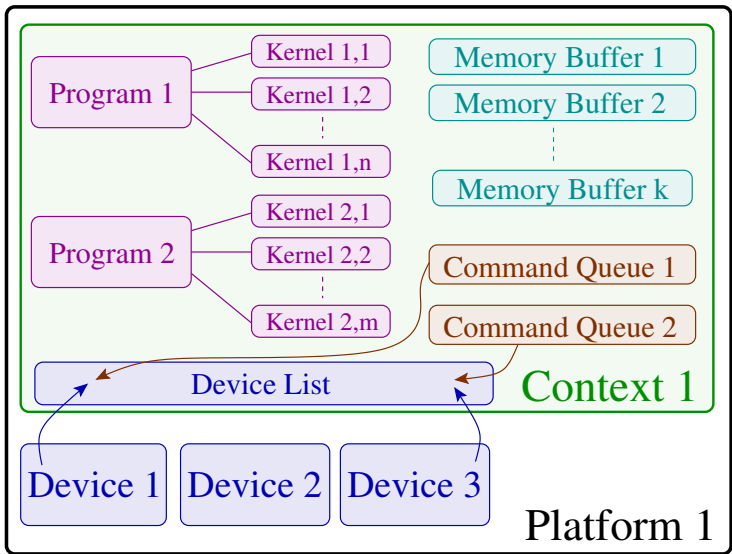
OpenCL Platform Model



OpenCL Platform Model



OpenCL Platform Model



OpenCL Thread Control (1D) vs. CUDA

- Local ID in block: `get_local_id(0)`
 - Threads per block: `get_local_size(0)`
 - ID of block: `get_group_id()`
 - No. of blocks: `get_num_groups()`
 - Global thread ID: `get_global_id()`
 - No. of threads: `get_global_size()`
 - Barrier: `barrier(CLK_GLOBAL_MEM_FENCE)`
- `threadIdx.x`
 - `blockDim.x`
 - `blockIdx.x`
 - `gridDim.x`
 -
 -
 - `__syncthreads()`

OpenCL Example

```
// Multiplies A*x, leaving the result in y.  
// A is a row-major matrix,  
// meaning the (i,j) element is at A[i*ncols+j].  
__kernel void matvec(__global const float *A,  
                    __global const float *x,  
                    uint ncols, __global float *y)  
{  
    size_t i = get_global_id(0);    // Global id, used as the row  
    index  
    __global float const *a = &A[i*ncols]; // Pointer to the i-th row  
    float sum = 0.f;                // Accumulator for dot product  
    for (size_t j = 0; j < ncols; j++) {  
        sum += a[j] * x[j];  
    }  
    y[i] = sum;  
}
```

Source: <https://en.wikipedia.org/wiki/OpenCL>

Exercises

Environment

- <https://gtx1080.360252.org/2020/ex7/>
- (Might receive visual updates and additional hints over the next days)
- Due: Tuesday, December 8, 2020 at 23:59pm

Hints and Suggestions

- Consider version control for locally developed code
- Please let me know of any bugs or issues

Freakin' Fast Friday

Opportunity for Informal Chatting

- When? Friday, December 4, 17:00-18:00
- Where? ~~Wieden Bräu~~ This Zoom channel
- What? Preserving mental sanity during Lockdown

Hints and Suggestions

- Consider bringing a drink
- Will not change your course evaluation
- **Completely optional and no obligation to show up**