



INSTITUTE OF MICROELECTRONICS

360.252 COMPUTATIONAL SCIENCE ON MANY-CORE ARCHITECTURES

## Exercise 7

Author:

*Peter* HOLZNER, 01426733

Submission: December 8, 2020

## Contents

<b>1 Task 1: Dot Product With OpenCL</b>	<b>1</b>
1.1 Implementation . . . . .	1
1.2 Runtime comparisons . . . . .	1
<b>2 Task 2: OpenCL Sparse Matrix-Vector Product Kernels</b>	<b>3</b>
<b>3 Code and Kernels</b>	<b>6</b>

---

## 1 Task 1: Dot Product With OpenCL

Code listings for this task:

- OpenCL dot product kernel: Listing 1
- OpenCL benchmark: Listing 2
- CUDA benchmark + kernel: Listing 3

### 1.1 Implementation

I've already aired my frustrations with this exercise in the chat. Considering Nvidias soft lock on OpenCL (still stuck on OCL v1.2), I thought it only to be fair to compare my OpenCL implementation to a similar CUDA kernel - also that's really the fairest benchmark anyway (same conditions for everyone). That means, that I had to rely on not using advanced synchronization constructs such warp shuffles and work group reduction (OCL >2.0) respectively. For OpenCL, I was also unable to use atomic reductions for the global reduction due to the Intel OpenCL platform not supporting the necessary 64bit atomic extensions. Remnants of my attempts (atomic\_add, different dot-kernel versions) can still be found in my code. That did turn into a bit of a rant anyway, sorry!

I decided to only do the first part of the dot product on the OCL device. The results of each workgroup are written to one entry of the result array. The final results is then calculated on the CPU via summing across all workgroup results. Since we have already written an identical benchmark for CUDA in Exercise 2, I simply reused that code (and fixed it, since mine did not work properly back then).

### 1.2 Runtime comparisons

In Fig. 1, the total runtimes (dot product on device + reduction across work groups/blocks on host) for the 3 different implementations are shown on the left side. The OpenCL implementations completely outperform the CUDA implementation for vectors of size  $N < 1e7$  - by multiple orders of magnitude. Ironically, the OCL-GPU kernel performs the best across the board up until the largest vector tested at least. I tried with larger  $N$  but the program simply exited without any messages - I assume that I simply wasn't able to allocate enough memory and did not want to cause any further problems with the device.

OpenCL does not seem to suffer from as much overhead as CUDA when starting a kernel, so OpenCL seems better to me than CUDA for smaller problem sizes and/or programs where many "smaller" kernels need to be called (for whatever reason). Apart from that, OpenCL performed better than CUDA for larger vectors as well, although I'm hesitant to make a definitive statement based on these limited benchmarks. OpenCL is much more complicated, unintuitive and annoying to

program for though - especially when one is bound to OCL v1.2 and needs to write code for multiple devices (though you can of course make your code more modular by using precompiler directives and enabling/disabling certain code parts). At least some of the aforementioned overhead that CUDA suffers from is probably hidden in some of the additional setup steps needed for the OpenCL kernel execution - both in terms of consideration needed from the programmer and in terms of the observed runtime.

The right side of Fig. 1 shows the breakdown of the two stages of the dot-product implementation. The final reduction on the host (CPU in both cases in plain C++) is clearly not the bottleneck - or at least not more than the kernel execution.

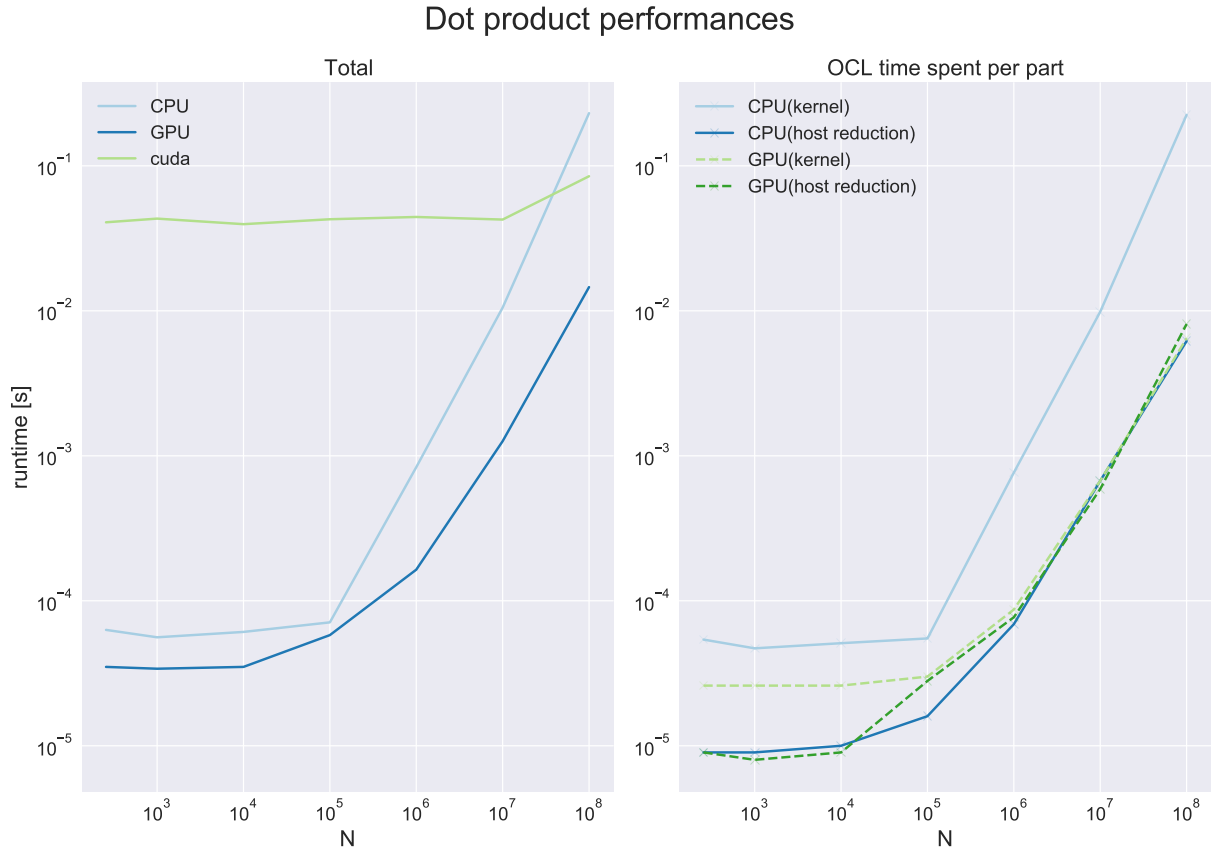


Fig. 1. Runtimes of the different vector dot product based kernels.

## 2 Task 2: OpenCL Sparse Matrix-Vector Product Kernels

Code listings for this task:

- OpenCL Matrix-Vector product kernel: Listing 4
- OpenCL Matrix-Vector product benchmark: Listing 5

Both kernels work and deliver the correct results as can be checked by running the benchmark code. The output looks as follows for both kernels:

### Run output

```
# Benchmarking finite difference matrix
Using the following device: GeForce GTX 1080
From OCL benchmark: 0.000331
Reference:
2 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | ... | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 2
OpenCL:
2 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | ... | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 2
Difference between the two: 0 (check: 1
Time for ref product: 1.6e-05
Time for OCL product: 0.000331
# Benchmarking special matrix
Using the following device: GeForce GTX 1080
From OCL benchmark: 0.001752
Reference:
-256 | -242 | -256 | -256 | -255 | -178 | -253 | -252 | -256 | -149 | ... | -256 | -256 | -229 | -233 | -255 | -255 | -251 | -228 | -252 | -247
OpenCL:
-256 | -242 | -256 | -256 | -255 | -178 | -253 | -252 | -256 | -149 | ... | -256 | -256 | -229 | -233 | -255 | -255 | -251 | -228 | -252 | -247
Difference between the two: 0 (check: 1
Time for ref product: 0.000219
Time for OCL product: 0.001752
Data:
Runtimes in csv form can be found here
https://gtx1080.360252.org/2020/ex7/ph\_data\_GPU.csv
```

Fig. 2. Sample output of the benchmark running the "slow" kernel.

My kernels did however not perform too well in terms of runtime, as can (unfortunately) be seen in Fig. 3. The GPU did not perform well for me in this case - it never beat the host computation. The OpenCL-CPU kernel did however outperform the reference host computation atleast for larger problem sizes or the denser matrix (with more nonzero entries per row).

## Sparse Matrix-Vector Product

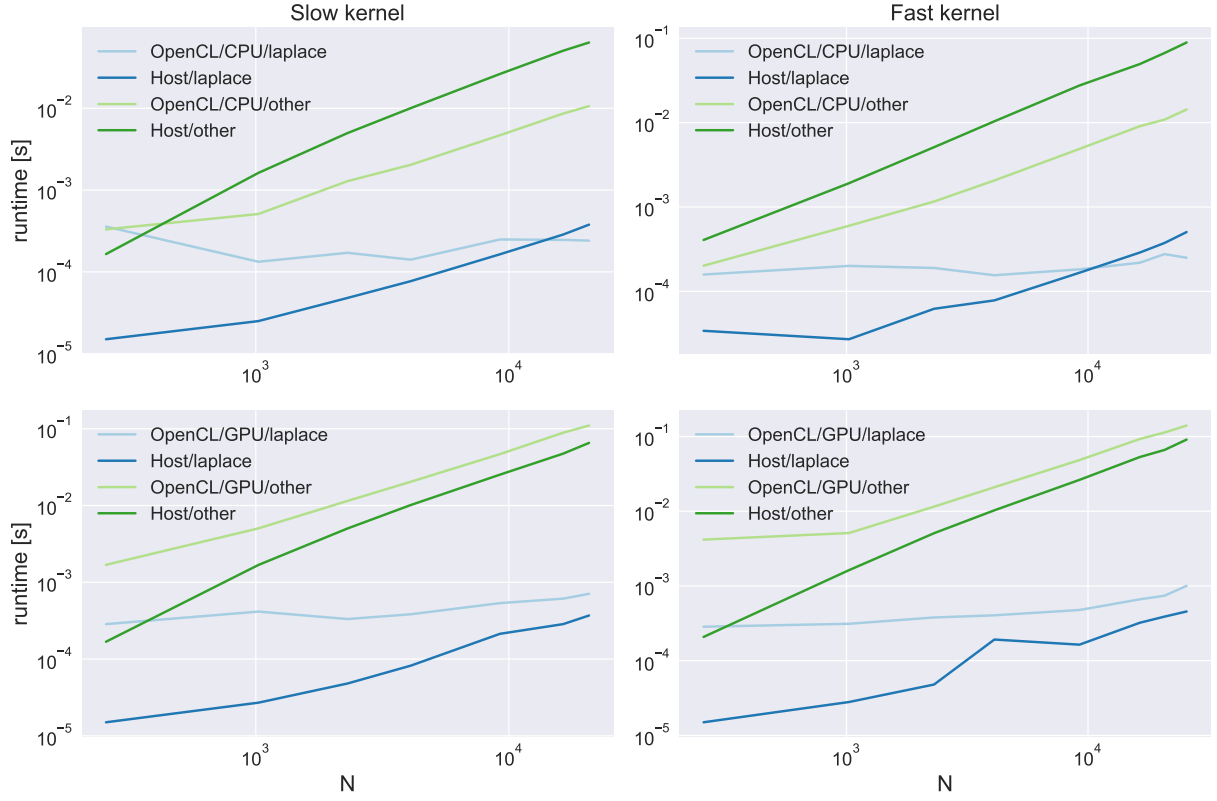


Fig. 3. Runtimes of the Sparse Matrix-Vector products. The left side shows the "slow" kernel, the right side the "faster" version. The matrix type is denoted by either "laplace" (blue lines) or "other" (green lines). The top two show the OpenCL-CPU runtimes, whereas the bottom shows the GPU runtimes.

In Fig. 4, the speedup of the "faster"<sup>1</sup> kernel version is plotted as compared to recorded runtime of the "slower" kernel. On the CPU, the faster kernel did actually improve the performance a little bit. On the GPU, I would classify it to be a tie.

<sup>1</sup>The slower kernel is called "ocl\_csr\_matvec". The faster kernel is called "ocl\_csr\_matvec\_fast". Very descriptive, I know.




Fig. 4. Speedup of the "faster" kernel compared to the "slower" kernel on the CPU and GPU respectively.

### 3 Code and Kernels

#### Listings

1	Ex7.1: Dot products with OpenCL - Kernels . . . . .	7
2	Ex7.1: Dot products with OpenCL - Benchmark . . . . .	8
3	Ex7.1: Dot products with CUDA . . . . .	14
4	Ex7.2: Sparse Matrix-Vector Product OpenCL - Kernels . . . . .	17
5	Ex7.2: Sparse Matrix-Vector Product OpenCL - Benchmark . . . . .	18

# Listing 1: Ex7.1: Dot products with OpenCL - Kernels

```

1  #pragma OPENCL EXTENSION cl_khr_fp64 : enable
2  #pragma OPENCL EXTENSION cl_khr_int64_base_atomics : enable
3  #pragma OPENCL EXTENSION cl_khr_int64_extended_atomics : enable
4
5  void my_atomic_add(volatile __global double *p, double val) {
6      volatile __global ulong* address_as_ul = (volatile __global ulong *) p;
7      volatile ulong old = *address_as_ul, assumed;
8      ulong val_as_ul = (ulong) val;
9      do {
10         assumed = old;
11         old = atomic_add(address_as_ul, val_as_ul)
12     } while (assumed != old);
13 };
14
15 __kernel void vec_add(__global double *x, __global double *y, unsigned int N) {
16     for (unsigned int i = get_global_id(0); i < N; i += get_global_size(0))
17         x[i] += y[i];
18 };
19
20 __kernel void xDOTy(__global double *result,
21                     __global double *x,
22                     __global double *y,
23                     __local cache, uint N) {
24     uint gid = get_global_id(0);
25     uint lid = get_local_id(0);
26     uint stride = get_global_size(0);
27     __local double cache[128];
28     double tmp = 0.0;
29     for (uint i = gid; i < N; i += stride)
30         tmp += x[i] * y[i];
31     cache[lid] = tmp;
32
33     for (int i = get_local_size(0) / 2; i > 0; i /= 2) {
34         barrier(CLK_LOCAL_MEM_FENCE);
35         if (lid < i)
36             cache[lid] += cache[lid + i];
37     }
38     if (lid==0)
39         result[get_group_id(0)] = cache[lid];
40 };
41
42
43
44 // __kernel void xDOTy(__global atomic_double *result, __global double *x,
45 //                     __global double *y, uint N) {
46 //     uint gid = get_global_id(0);
47 //     uint lid = get_local_id(0);
48 //     uint stride = get_global_size(0);
49 //     double dot = 0.0;
50 //     for (uint i = tid; i < N; i += stride)
51 //         dot += x[i] * y[i];
52 //
53 //     // need to add to signature: __local double *cache,
54 //     // for (int i = get_local_size(0) / 2; i > 0; i /= 2) {
55 //     //     barrier(CLK_LOCAL_MEM_FENCE);
56 //     //     if (lid < i)
57 //     //         cache[lid] += cache[lid + i];
58 //     // }
59 //     double val = work_group_reduce_add(gid < N ? dot : 0.);
60 //
61 //     if (lid == 0) {
62 //         atomic_add(result, val);
63 //     }
64 //     // // atomic flag version
65 //     // // need to add: __global volatile atomic_flag *lock,
66 //     // // if (lid == 0) {
67 //     //     while (!atomic_flag_test_and_set(lock)){
68 //     //         *result += cache[0];
69 //     //         atomic_flag_clear(lock);
70 //     //     }

```



```

71 // };
72
73 // __kernel void xDOTy(__global atomic_double *result, __global double *x,
74 //                     __global double *y, __local double *cache,
75 //                     __global bool *lock, uint N) {
76 //     uint gid = get_global_id(0);
77 //     uint lid = get_local_id(0);
78 //     uint stride = get_global_size(0);
79 //     double dot = 0.0;
80 //     for (uint i = tid; i < N; i += stride)
81 //         dot += x[i] * y[i];
82 //     #cache[lid] = dot;
83
84 //     // for (int i = get_local_size(0) / 2; i > 0; i /= 2) {
85 //         barrier(CLK_LOCAL_MEM_FENCE);
86 //         if (lid < i)
87 //             cache[lid] += cache[lid + i];
88 //     }
89 //     double val = work_group_reduce_max( gid < length ? input[globalId] : -INFINITY);
90
91 //     if (lid == 0) {
92 //         while (*lock) {}
93
94 //         *result += cache[0];
95 //     }
96 // };

```

Listing 2: Ex7.1: Dot products with OpenCL - Benchmark

```

1 //
2 // Tutorial for demonstrating a simple OpenCL vector addition kernel
3 //
4 // Author: Karl Rupp      rupp@iue.tuwien.ac.at
5 //
6 typedef double ScalarType;
7 #include <iostream>
8 #include <fstream>
9 #include <string>
10 #include <vector>
11 #include <cmath>
12 #include <algorithm>
13 #include <numeric>
14 #include <stdexcept>
15
16 #ifdef __APPLE__
17 #include <OpenCL/cl.h>
18 #else
19 #include <CL/cl.h>
20 #endif
21
22 // Helper include file for error checking
23 #include "ocl-error.hpp"
24 #include "timer.hpp"
25
26 #define LOCAL_SIZE 128
27 #define GLOBAL_SIZE 128
28 #define NUM_TESTS 5
29
30 template <typename T>
31 void printContainer(T container, const int size) {
32     std::cout << container[0];
33     for (int i = 1; i < size; ++i)
34         std::cout << " | " << container[i] ;
35     std::cout << std::endl;
36 }
37
38 template <typename T>
39 void printContainer(T container, const int size, const int only) {
40     std::cout << container[0];
41     for (int i = 1; i < only; ++i)
42         std::cout << " | " << container[i];
43     std::cout << " | ...";

```

```

44     for (int i = size - only; i < size; ++i)
45         std::cout << " | " << container[i];
46     std::cout << std::endl;
47 }
48
49 double median(std::vector<double>& vec)
50 {
51     size_t size = vec.size();
52     if (size == 0)
53         return 0.;
54     sort(vec.begin(), vec.end());
55     size_t mid = size/2;
56
57     return size % 2 == 0 ? (vec[mid] + vec[mid-1]) / 2 : vec[mid];
58 }
59
60 // const char *my_ompcl_program = "\n"
61 // "#pragma OPENCL EXTENSION cl_khr_fp64 : enable\n"
62 // " \n\n"
63 // "__kernel void xDOTy(__global double *result, \n"
64 // "                        __global double *x,\n"
65 // "                        __global double *y, \n"
66 // "                        __local cache, uint N) {\n"
67 // "    uint gid = get_global_id(0);\n"
68 // "    uint lid = get_local_id(0);\n"
69 // "    uint stride = get_global_size(0);\n"
70 // "    double tmp = 0.0;\n"
71 // "    for (uint i = gid; i < N; i += stride)\n"
72 // "        tmp = x[i] * y[i];\n"
73 // "    cache[lid] = tmp;\n"
74 // "    \n"
75 // "    for (int i = get_local_size(0) / 2; i > 0; i /= 2) {\n"
76 // "        barrier(CLK_LOCAL_MEM_FENCE);\n"
77 // "        if (lid < i)\n"
78 // "            cache[lid] += cache[lid + i];\n"
79 // "    }\n"
80 // "    if (lid==0)\n"
81 // "        result[get_group_id()] = cache[lid];\n"
82 // "};\n";
83
84
85 // "                        __local cache, uint N) {\n"
86
87 std::string my_ompcl_program_c = "\n"
88 "#pragma OPENCL EXTENSION cl_khr_fp64 : enable\n"
89 " \n\n"
90 "__kernel void xDOTy(__global double *result, \n"
91 "                    __global double *x,\n"
92 "                    __global double *y, \n"
93 "                    uint N) {\n"
94 "    uint gid = get_global_id(0);\n"
95 "    uint lid = get_local_id(0);\n"
96 "    uint stride = get_global_size(0);\n"
97 "    __local double cache[128];\n"
98 "    double tmp = 0.0;\n"
99 "    for (uint i = gid; i < N; i += stride)\n"
100 "        tmp += x[i] * y[i];\n"
101 "    cache[lid] = tmp;\n"
102 "    \n"
103 "    for (int i = get_local_size(0) / 2; i > 0; i /= 2) {\n"
104 "        barrier(CLK_LOCAL_MEM_FENCE);\n"
105 "        if (lid < i)\n"
106 "            cache[lid] += cache[lid + i];\n"
107 "    }\n"
108 "    if (lid==0)\n"
109 "        result[get_group_id(0)] = cache[lid];\n"
110 "};\n";
111
112 std::string header = "\n"
113 "#pragma OPENCL EXTENSION cl_khr_fp64 : enable\n"
114 " \n\n";
115 std::string kernel_signature = "__kernel void xDOTy";

```

```

116 std::string kernel_code1 = "(__global double *result, \n"
117 "    __global double *x,\n"
118 "    __global double *y, \n"
119 "    uint N) {\n";
120 std::string kernel_code2 = "    uint gid = get_global_id(0);\n"
121 "    uint lid = get_local_id(0);\n"
122 "    uint stride = get_global_size(0);\n"
123 "    __local double cache[128];\n"
124 "    double tmp = 0.0;\n"
125 "    for (uint i = gid; i < N; i += stride)\n"
126 "        tmp += x[i] * y[i];\n"
127 "    cache[lid] = tmp;\n"
128 "    \n"
129 "    for (int i = get_local_size(0) / 2; i > 0; i /= 2) {\n"
130 "        barrier(CLK_LOCAL_MEM_FENCE);\n"
131 "        if (lid < i)\n"
132 "            cache[lid] += cache[lid + i];\n"
133 "    }\n"
134 "    if (lid==0)\n"
135 "        result[get_group_id(0)] = cache[lid];\n"
136 "};\n";
137
138
139 int main()
140 {
141     cl_int err;
142
143     bool compute = false;
144     bool compile_M = true;
145     std::vector<uint> M_vec{1, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100};
146     if (!compile_M) {
147         M_vec.clear();
148         M_vec.push_back(1);
149     }
150     // std::string target = "GeForce GTX 1080";
151     std::string target = "GPU";
152     std::vector<uint> N_vec{256, 1000, 10000, 100000, 1000000, 10000000, 100000000};
153     uint N_min = N_vec.front();
154     uint N_max = N_vec.back();
155     uint cnt = N_vec.size();
156
157
158     Timer timer;
159     timer.reset();
160
161     //
162     ////////////////////////////////////////////////// Part 1: Set up an OpenCL context with one device
163     //////////////////////////////////////////////////
164     //
165     //
166     // Query platform:
167     //
168     cl_uint num_platforms;
169     cl_platform_id platform_ids[42]; //no more than 42 platforms supported...
170     err = clGetPlatformIDs(42, platform_ids, &num_platforms); OPENCL_ERR_CHECK(err);
171     std::cout << "# Platforms found: " << num_platforms << std::endl;
172
173
174     //
175     // Query devices:
176     //
177     cl_device_id device_ids[42];
178     cl_uint num_devices;
179     char device_name[64];
180     cl_device_id my_device_id;
181     cl_platform_id my_platform;
182     for (int i = 0; i < num_platforms; ++i)
183     {
184         my_platform = platform_ids[i];
185         if (target == "GPU") {
186             err = clGetDeviceIDs(my_platform, CL_DEVICE_TYPE_GPU, 42, device_ids, &num_devices);

```

```

187     }
188     else {
189         err = clGetDeviceIDs(my_platform, CL_DEVICE_TYPE_CPU, 42, device_ids, &num_devices);
190     }
191     if (err == CL_SUCCESS)
192         break;
193 }
194 OPENCL_ERR_CHECK(err);
195 std::cout << "# Devices found: " << num_devices << std::endl;
196 my_device_id = device_ids[0];
197
198 size_t device_name_len = 0;
199 err = clGetDeviceInfo(my_device_id, CL_DEVICE_NAME, sizeof(char)*63, device_name, &
    device_name_len); OPENCL_ERR_CHECK(err);
200
201 std::cout << "Using the following device: " << device_name << std::endl;
202
203 //
204 // Create context:
205 //
206 cl_context my_context = clCreateContext(0, 1, &my_device_id, NULL, NULL, &err);
    OPENCL_ERR_CHECK(err);
207
208
209 //
210 // create a command queue for the device:
211 //
212 cl_command_queue my_queue = clCreateCommandQueueWithProperties(my_context, my_device_id,
    0, &err); OPENCL_ERR_CHECK(err);
213
214
215 //
216 ////////////////////////////////////////////////// Part 2: Create a program and extract kernels
    //////////////////////////////////////////////////
217 //
218 //
219 // Build the program:
220 //
221 cl_kernel my_kernel;
222 cl_program prog;
223
224
225 std::fstream csv_compile;
226 std::string csv_compile_name = "ph_data_compile.csv";
227
228 if (compile_M) {
229     csv_compile.open(csv_compile_name, std::fstream::out | std::fstream::trunc);
230     csv_compile << "M;compile_time;create_time" << std::endl;
231 }
232 int m = 1;
233
234 std::string ocl_prog = header + kernel_signature + kernel_code1 + kernel_code2;
235 for (auto& M : M_vec){
236
237     // // To generate the M kernels
238     for (; m < M; ++m) {
239         ocl_prog += kernel_signature + std::to_string(m) + kernel_code1 + "uint insert" + std
            ::to_string(m) + "=1;\n" + kernel_code2;
240     }
241     const char * my_opencl_program = ocl_prog.c_str();
242
243     std::vector<double> tmp(NUM_TESTS, 0);
244     for (uint iter = 0; iter < NUM_TESTS; iter++){
245         timer.reset();
246         size_t source_len = std::string(my_opencl_program).length();
247         prog = clCreateProgramWithSource(my_context, 1, &my_opencl_program, &source_len, &err)
            ; OPENCL_ERR_CHECK(err);
248         err = clBuildProgram(prog, 0, NULL, NULL, NULL, NULL);
249         tmp[iter] = timer.get();
250     }
251     double compile_time = median(tmp);
252     //

```

```

253 // Print compiler errors if there was a problem:
254 //
255 if (err != CL_SUCCESS) {
256
257     char *build_log;
258     size_t ret_val_size;
259     err = clGetProgramBuildInfo(prog, my_device_id, CL_PROGRAM_BUILD_LOG, 0, NULL, &
        ret_val_size);
260     build_log = (char *)malloc(sizeof(char) * (ret_val_size+1));
261     err = clGetProgramBuildInfo(prog, my_device_id, CL_PROGRAM_BUILD_LOG, ret_val_size,
        build_log, NULL);
262     build_log[ret_val_size] = '\0'; // terminate string
263     std::cout << "Log: " << build_log << std::endl;
264     free(build_log);
265     std::cout << "OpenCL program sources: " << std::endl << my_opencl_program << std::endl
        ;
266     return EXIT_FAILURE;
267 }
268 //
269 // Extract the only kernel in the program:
270 //
271 for (uint iter = 0; iter < NUM_TESTS; iter++){
272     timer.reset();
273     my_kernel = clCreateKernel(prog, "xDOTy", &err); OPENCL_ERR_CHECK(err);
274     tmp[iter] = timer.get();
275 }
276 double create_time = median(tmp);
277 std::cout << "Time to compile and create kernel: " << compile_time << std::endl;
278 csv_compile << M << ";";
279     << compile_time << ";";
280     << create_time << std::endl;
281 }
282
283
284 // Plan to reuse all these vectors and buffers
285 double y_val = 2., x_val=1.;
286 std::vector<ScalarType> x(N_max, x_val);
287 std::vector<ScalarType> y(N_max, y_val);
288 std::vector<ScalarType> dot_vec(N_max, 0);
289
290 cl_mem ocl_x = clCreateBuffer(my_context, CL_MEM_READ_WRITE | CL_MEM_COPY_HOST_PTR, N_max
    * sizeof(ScalarType), &(x[0]), &err); OPENCL_ERR_CHECK(err);
291 cl_mem ocl_y = clCreateBuffer(my_context, CL_MEM_READ_WRITE | CL_MEM_COPY_HOST_PTR, N_max
    * sizeof(ScalarType), &(y[0]), &err); OPENCL_ERR_CHECK(err);
292 cl_mem ocl_dot = clCreateBuffer(my_context, CL_MEM_READ_WRITE | CL_MEM_COPY_HOST_PTR,
    sizeof(ScalarType)*N_max, dot_vec.data(), &err); OPENCL_ERR_CHECK(err);
293
294 //
295 ////////////////////////////////////////////////// Part 3: Create memory buffers
    //////////////////////////////////////
296 //
297
298 std::string csv_name = "ph_data_" + target + ".csv";
299 std::fstream csv;
300 if (compute){
301     std::vector<double> times_cpu(N_vec.size(), 0.);
302     std::vector<double> times_gpu(N_vec.size(), 0.);
303     std::vector<double> times_total(N_vec.size(), 0.);
304     csv.open(csv_name, std::fstream::out | std::fstream::trunc);
305     csv << "N;target;local_size;global_size;ocl_time;cpu_time;total_time;dot" << std::endl;
306     for (auto& N: N_vec){
307
308         cl_uint vector_size = N;
309         size_t local_size = LOCAL_SIZE;
310         size_t global_size = GLOBAL_SIZE*GLOBAL_SIZE;
311         size_t groups = 1 + int(N/LOCAL_SIZE);
312
313         //
314         ////////////////////////////////// Part 4: Run kernel //////////////////////////////////
315         //
316
317

```

```

318 //
319 // Set kernel arguments:
320 //
321 // xDOTy(__global double *result,
322 //        __global double *x,
323 //        __global double *y,
324 //        __local cache, uint N)
325 err = clSetKernelArg(my_kernel, 0, sizeof(cl_mem), (double*)&ocl_dot);
326     OPENCL_ERR_CHECK(err);
327 err = clSetKernelArg(my_kernel, 1, sizeof(cl_mem), (double*)&ocl_x); OPENCL_ERR_CHECK
328     (err);
329 err = clSetKernelArg(my_kernel, 2, sizeof(cl_mem), (double*)&ocl_y); OPENCL_ERR_CHECK
330     (err);
331 // err = clSetKernelArg(my_kernel, 3, sizeof(cl_float) * local_work_size[0], NULL);
332     OPENCL_ERR_CHECK(err);
333 err = clSetKernelArg(my_kernel, 3, sizeof(cl_uint), (void*)&vector_size);
334     OPENCL_ERR_CHECK(err);
335
336 //
337 // Enqueue kernel in command queue:
338 //
339 std::vector<double> tmp(NUM_TESTS, 0);
340 for (uint iter = 0; iter < NUM_TESTS; iter++){
341     timer.reset();
342     err = clEnqueueNDRangeKernel(my_queue, my_kernel, 1, NULL, &global_size, &local_size
343         , 0, NULL, NULL); OPENCL_ERR_CHECK(err);
344
345     // wait for all operations in queue to finish:
346     err = clFinish(my_queue); OPENCL_ERR_CHECK(err);
347     tmp[iter] = timer.get();
348 }
349 double ocl_time = median(tmp);
350
351 //
352 ////////////////////////////////////////////////// Part 5: Get data from OpenCL buffer
353 //////////////////////////////////////////////////
354 //
355 // err = clEnqueueReadBuffer(my_queue, ocl_x, CL_TRUE, 0, sizeof(ScalarType) * x.size
356 //     (), &(x[0]), 0, NULL, NULL); OPENCL_ERR_CHECK(err);
357 err = clEnqueueReadBuffer(my_queue, ocl_dot, CL_TRUE, 0, sizeof(ScalarType)*
358     vector_size, dot_vec.data(), 0, NULL, NULL); OPENCL_ERR_CHECK(err);
359
360 timer.reset();
361 double dot = std::accumulate(dot_vec.begin(), dot_vec.begin()+groups, 0.);
362 double cpu_time = timer.get();
363 double total_time = ocl_time + cpu_time;
364
365 std::cout << std::endl;
366 std::cout << "Result of kernel execution: " << dot;
367 std::cout << " =? " << vector_size*y_val*x_val << " : " << (dot == y_val*x_val*
368     vector_size) << std::endl;
369 std::cout << "Runtime: kernel + cpu = total_time" << std::endl;
370 std::cout << ocl_time << " + " << cpu_time << " = " << total_time << std::endl;
371
372 csv << N<< ";";
373     << target << ";";
374     << local_size<< ";";
375     << global_size<< ";";
376     << ocl_time<< ";";
377     << cpu_time<< ";";
378     << total_time << ";";
379     << dot << std::endl;
380
381 // std::cout << "Result container:" << std::endl;
382 // printContainer(dot_vec, dot_vec.size(), 10);
383 }
384 csv.close();
385 }
386 if (compile_M) csv_compile.close();

```

```

380 //
381 // cleanup
382 //
383 clReleaseMemObject(ocl_x);
384 clReleaseMemObject(ocl_y);
385 clReleaseMemObject(ocl_dot);
386 clReleaseProgram(prog);
387 clReleaseCommandQueue(my_queue);
388 clReleaseContext(my_context);
389
390 std::cout << "Data: \nRuntimes in csv form can be found here\nhttps://gtx1080.360252.org
/2020/ex7/" << csv_name << std::endl;
391 std::cout << "Data: \nCompile times in csv form can be found here\nhttps://gtx1080.360252.
org/2020/ex7/" << csv_compile_name << std::endl;
392
393 std::cout << std::endl;
394 std::cout << "#" << std::endl;
395 std::cout << "# My first OpenCL application finished successfully!" << std::endl;
396 std::cout << "#" << std::endl;
397
398 std::cout << "And here it is:" << std::endl;
399 std::cout << ocl_prog;
400 return EXIT_SUCCESS;
401 }

```

Listing 3: Ex7.1: Dot products with CUDA

```

1  #include <iostream>
2  #include <string>
3  #include <vector>
4  #include "timer.hpp"
5  #include <cmath>
6
7  #define BLOCK_SIZE 256
8
9  __global__ void initKernel1(double* arr, const double value, const size_t N)
10 {
11     const int stride = blockDim.x * gridDim.x;
12     int tid = threadIdx.x + blockIdx.x * blockDim.x;
13
14     for(; tid < N; tid += stride)
15     {
16         arr[tid] = value;
17     }
18 }
19
20 __global__ void initKernel2(double* arr, double* arr2, const size_t N)
21 {
22     const int stride = blockDim.x * gridDim.x;
23     int tid = threadIdx.x + blockIdx.x * blockDim.x;
24
25     for(; tid < N; tid += stride)
26     {
27         arr[tid] = tid;
28         arr2[tid] = N - 1 - tid;
29     }
30 }
31
32 __global__ void initKernel3(double* arr, double* arr2, double* arr3, const size_t N)
33 {
34     const int stride = blockDim.x * gridDim.x;
35     int tid = threadIdx.x + blockIdx.x * blockDim.x;
36
37     for(; tid < N; tid += stride)
38     {
39         arr[tid] = tid;
40         arr2[tid] = N - 1 - tid;
41         arr3[tid] = 0;
42     }
43 }
44
45 __global__ void addKernel(double* x, double* y, double* res, const size_t N)

```

```

46 {
47     const int stride = blockDim.x * gridDim.x;
48     int tid = threadIdx.x + blockIdx.x * blockDim.x;
49
50     for(; tid < N; tid += stride)
51     {
52         res[tid] = x[tid] + y[tid];
53         //res[tid] += 1;
54     }
55 }
56
57 __global__ void dot_A_1(double* x, double* y, double* block_sums, const size_t N)
58 {
59     uint tid = threadIdx.x + blockDim.x * blockIdx.x;
60     uint stride = blockDim.x * gridDim.x;
61
62     __shared__ double cache[BLOCK_SIZE];
63
64     double tid_sum = 0.0;
65     for (; tid < N; tid += stride)
66     {
67         tid_sum += x[tid] * y[tid];
68     }
69     cache[threadIdx.x] = tid_sum;
70
71     __syncthreads();
72     for (uint i = blockDim.x/2; i != 0; i /=2)
73     {
74         if (threadIdx.x < i) //lower half does smth, rest idles
75         {
76             cache[threadIdx.x] += cache[threadIdx.x + i]; //lower looks up by stride and sums up
77         }
78         __syncthreads();
79     }
80
81     if(threadIdx.x == 0) // cache[0] now contains block_sum
82     {
83         block_sums[blockIdx.x] = cache[0];
84     }
85     __syncthreads();
86 }
87
88 __global__ void dot_A_2(double* block_sums)
89 {
90     int tid = threadIdx.x; // only one block, so this is fine!
91     for (int i = blockDim.x / 2; i >= 1; i /=2) // same principal as above
92     {
93         if (tid < i)
94         {
95             block_sums[tid] += block_sums[tid + i];
96         }
97         __syncthreads();
98     }
99 }
100
101 __global__ void dot_Atomic(double* x, double* y, double* result, const size_t N)
102 {
103     uint tid = threadIdx.x + blockDim.x * blockIdx.x ;
104     uint stride = blockDim.x * gridDim.x ;
105
106     __shared__ double cache[BLOCK_SIZE];
107
108     double tid_sum = 0.0;
109     for (; tid < N; tid += stride)
110     {
111         tid_sum += x[tid] * y[tid];
112     }
113     tid = threadIdx.x;
114     cache[tid] = tid_sum;
115
116     __syncthreads();
117     for (uint i = blockDim.x/2; i != 0; i /=2)

```



```

118 {
119     __syncthreads();
120     if (tid < i) //lower half does smth, rest idles
121     {
122         cache[tid] += cache[tid + i ]; //lower looks up by stride and sums up
123     }
124 }
125
126 if(tid == 0) // cache[0] now contains block_sum
127 {
128     atomicAdd(result, cache[0]);
129 }
130 }
131
132 int main(void)
133 {
134     const int num_tests = 10;
135     int tests_done = num_tests;
136     const int block_size = BLOCK_SIZE;
137     const int mode = 1;
138     const int option = 2;
139
140     std::cout << "N;blocks;block_size;tests_done;total_time;time_per_test;check" << std::endl;
141     for (size_t N = 256; N <= 1000000000; N*=10)
142     {
143         std::cout << N << std::endl;
144         //int blocks = (int)(N+block_size-1)/block_size;
145         int blocks = block_size;
146         double result = 0.0;
147         double result_true = N;
148         double* h_block_sums = (double *)malloc(sizeof(double) * blocks);
149         double* presult = (double *)malloc(sizeof(double));
150         presult = &result;
151         double* pnull = (double *)malloc(sizeof(double));
152         *pnull = 0.0;
153         double *d_x, *d_y, *d_block_sums, *d_result;
154
155         cudaMalloc(&d_result, sizeof(double));
156         cudaMalloc(&d_x, N*sizeof(double));
157         cudaMalloc(&d_y, N*sizeof(double));
158         cudaMalloc(&d_block_sums, blocks*sizeof(double));
159         cudaDeviceSynchronize();
160
161         int i = 0;
162         Timer timer;
163         double total_time = 0.0;
164
165         initKernel1<<<blocks, block_size>>>(d_block_sums, 0.0, block_size);
166         initKernel1<<<blocks, block_size>>>(d_x, 1.0, N);
167         initKernel1<<<blocks, block_size>>>(d_y, 1.0, N);
168         //initKernel2<<<blocks, block_size>>>(d_x, d_y, N);
169         cudaDeviceSynchronize();
170
171         timer.reset();
172         for (i = 0; i<num_tests; i++)
173         {
174             tests_done = i+1;
175             if (option == 1)
176             {
177                 dot_A_1<<<blocks, block_size>>>(d_x, d_y, d_block_sums, N);
178                 dot_A_2<<<1, block_size>>>(d_block_sums);
179                 cudaDeviceSynchronize();
180             }
181             if (option == 2)
182             {
183                 dot_A_1<<<blocks, block_size>>>(d_x, d_y, d_block_sums, N);
184                 cudaMemcpy(h_block_sums, d_block_sums, blocks*sizeof(double), cudaMemcpyDeviceToHost);
185                 //std::cout << h_block_sums[0] << "=? " << h_block_sums[blocks-1] << std::endl;
186                 cudaDeviceSynchronize();
187                 result = 0.0;
188                 for (int j = 0; j < blocks; j+=1)

```

```

189         {
190             result += h_block_sums[j];
191         }
192     }
193     if (option == 3)
194     {
195         cudaMemcpy(d_result, pnull, sizeof(double), cudaMemcpyHostToDevice);
196         dot_Atomic<<<blocks, block_size>>>(d_x, d_y, d_result, N);
197     }
198
199     //std::cout << "(" << i+1 << ") Elapsed: " << runtime << " s" << std::endl;
200 }
201 total_time = timer.get();
202 size_t check;
203 if (option == 1)
204 {
205     cudaMemcpy(presult, d_block_sums, sizeof(double), cudaMemcpyDeviceToHost);
206     check = result - result_true;
207 }
208 if (option == 2)
209 {
210     check = result - result_true;
211 }
212 if (option == 3)
213 {
214     cudaMemcpy(presult, d_result, sizeof(double), cudaMemcpyDeviceToHost);
215     check = result - result_true;
216 }
217
218 if (mode == 0)
219 {
220     std::cout << std::endl << "Results after " << tests_done << " tests:" << std::endl;
221     std::cout << "Total runtime: " << total_time << std::endl;
222     std::cout << "Average runtime; " << total_time/tests_done << std::endl;
223     std::cout << "Check: " << result << " ?= " << result_true << std::endl;
224     std::cout << "\n\n";
225 }
226 if (mode == 1)
227 {
228     std::cout << N << ";"
229     << blocks << ";"
230     << block_size << ";"
231     << tests_done << ";"
232     << total_time << ";"
233     << total_time/tests_done << ";"
234     << check << std::endl;
235 }
236 cudaFree(d_x);
237 cudaFree(d_y);
238 cudaFree(d_block_sums);
239 cudaFree(d_result);
240 // free(presult);
241 // free(pnull);
242 free(h_block_sums);
243 if (N==256)
244 {
245     N=100;
246 }
247 }
248 cudaDeviceSynchronize();
249 return EXIT_SUCCESS;
250 }

```

Listing 4: Ex7.2: Sparse Matrix-Vector Product OpenCL - Kernels

```

1  #pragma OPENCL EXTENSION cl_khr_fp64 : enable
2
3  __kernel void ocl_csr_matvec(uint N,
4                               __global int *csr_rowoffsets,
5                               __global int *csr_colindices,
6                               __global double *csr_values,
7                               __global double const *x, __global double *y)

```

```

8  {
9      uint gid = get_global_id(0);
10     uint stride = get_global_size(0);
11     for (size_t i=gid; i<N; i+=stride) {
12         double value = 0;
13         for (size_t j=csr_rowoffsets[i]; j<csr_rowoffsets[i+1]; ++j)
14             value += csr_values[j] * x[csr_colindices[j]];
15         y[i] = value;
16     }
17 };
18
19 __kernel void ocl_csr_matvec_fast(uint N,
20                                 __global int *csr_rowoffsets,
21                                 __global int *csr_colindices,
22                                 __global double *csr_values,
23                                 __global double const *x, __global double *y)
24 {
25     __local double cache[128];
26     uint gid = get_group_id(0);
27     uint lid = get_local_id(0);
28     uint stride = get_num_groups(0);
29     uint j_stride = get_global_size(0);
30     for (size_t i=gid; i<N; i+=stride) {
31         double value = 0;
32         for (size_t j=csr_rowoffsets[i]; j<csr_rowoffsets[i+1]; j+=j_stride)
33             value += csr_values[j] * x[csr_colindices[j]];
34         cache[lid] = value;
35         for (int i = get_local_size(0) / 2; i > 0; i /= 2) {
36             barrier(CLK_LOCAL_MEM_FENCE);
37             if (lid < i)
38                 cache[lid] += cache[lid + i];
39         }
40         if (lid==0)
41             y[gid] = cache[0];
42     }
43 };

```

Listing 5: Ex7.2: Sparse Matrix-Vector Product OpenCL - Benchmark

```

1  #include <iostream>
2  #include <fstream>
3  #include <string>
4  #include <vector>
5  #include <cmath>
6  #include <algorithm>
7  #include <numeric>
8  #include <stdexcept>
9  #include "generate.hpp"
10 #include "timer.hpp"
11
12 #ifdef __APPLE__
13 #include <OpenCL/cl.h>
14 #else
15 #include <CL/cl.h>
16 #endif
17
18 // Helper include file for error checking
19 #include "ocl-error.hpp"
20
21 // #ifndef uint
22 // // not defined for me...it's annoying
23 // using uint = uint32_t;
24 // #endif
25 typedef double ScalarType;
26
27 #define LOCAL_SIZE 128
28 #define GLOBAL_SIZE 128
29 #define NUM_TESTS 5
30 #define TARGET "GPU"
31 #define TRUNC_CSV
32 #define PP 16
33

```

```

34 #define SLOW_KERNEL
35 #ifndef SLOW_KERNEL
36 #define FAST_KERNEL
37 #endif
38 std::string target = TARGET;
39 #ifdef SLOW_KERNEL
40 std::string csv_name = "ph_data_" + target + ".csv";
41 #endif
42 #ifdef FAST_KERNEL
43 std::string csv_name = "ph_data2_" + target + ".csv";
44 #endif
45
46 std::string my_opencl_program = "\n"
47 "#pragma OPENCL EXTENSION cl_khr_fp64 : enable\n"
48 "\n\n"
49 "__kernel void ocl_csr_matvec(uint N,\n"
50 "    __global int *csr_rowoffsets,\n"
51 "    __global int *csr_colindices,\n"
52 "    __global double *csr_values,\n"
53 "    __global double const *x, __global double *y)\n"
54 "{\n"
55 "    uint gid = get_global_id(0);\n"
56 "    uint stride = get_global_size(0);\n"
57 "    for (size_t i=gid; i<N; i+=stride) {\n"
58 "        double value = 0;\n"
59 "        for (size_t j=csr_rowoffsets[i]; j<csr_rowoffsets[i+1]; ++j)\n"
60 "            value += csr_values[j] * x[csr_colindices[j]];\n"
61 "        y[i] = value;\n"
62 "    }\n"
63 "};\n"
64 "\n\n"
65 "__kernel void ocl_csr_matvec_fast(uint N,\n"
66 "    __global int *csr_rowoffsets,\n"
67 "    __global int *csr_colindices,\n"
68 "    __global double *csr_values,\n"
69 "    __global double const *x, __global double *y)\n"
70 "{\n"
71 "    __local double cache[128];\n"
72 "    uint gid = get_group_id(0);\n"
73 "    uint lid = get_local_id(0);\n"
74 "    uint stride = get_num_groups(0);\n"
75 "    uint j_stride = get_global_size(0);\n"
76 "    for (size_t i=gid; i<N; i+=stride) {\n"
77 "        double value = 0;\n"
78 "        for (size_t j=csr_rowoffsets[i]; j<csr_rowoffsets[i+1]; j+=j_stride)\n"
79 "            value += csr_values[j] * x[csr_colindices[j]];\n"
80 "        cache[lid] = value;\n"
81 "        for (int i = get_local_size(0) / 2; i > 0; i /= 2) {\n"
82 "            barrier(CLK_LOCAL_MEM_FENCE);\n"
83 "            if (lid < i)\n"
84 "                cache[lid] += cache[lid + i];\n"
85 "        }\n"
86 "        if (lid==0)\n"
87 "            y[gid] = cache[0];\n"
88 "    }\n"
89 "};\n";
90
91 template <typename T>
92 void printContainer(T container, const int size) {
93     std::cout << container[0];
94     for (int i = 1; i < size; ++i)
95         std::cout << " | " << container[i] ;
96     std::cout << std::endl;
97 }
98
99 template <typename T>
100 void printContainer(T container, const int size, const int only) {
101     std::cout << container[0];
102     for (int i = 1; i < only; ++i)
103         std::cout << " | " << container[i];
104     std::cout << " | ...";
105     for (int i = size - only; i < size; ++i)

```

```

106     std::cout << " | " << container[i];
107     std::cout << std::endl;
108 }
109
110 double median(std::vector<double>& vec)
111 {
112     size_t size = vec.size();
113     if (size == 0)
114         return 0.;
115     sort(vec.begin(), vec.end());
116     size_t mid = size/2;
117
118     return size % 2 == 0 ? (vec[mid] + vec[mid-1]) / 2 : vec[mid];
119 };
120
121 bool check(const double* test, const double* ref, const uint N) {
122     for (uint i = 0; i < N; ++i){
123         if (test[i] != ref[i])
124             return false;
125     }
126     return true;
127 }
128
129 double diff_norm(const double* test, const double* ref, const uint N) {
130     double norm = 0.0;
131     for (uint i = 0; i < N; ++i){
132         norm += test[i] != ref[i];
133     }
134     return sqrt(norm);
135 }
136
137
138 /** Computes  $y = Ax$  for a sparse matrix A in CSR format and vector x,y. CPU implementation.
139 */
140 void csr_matvec_product(size_t N,
141                        int *csr_rowoffsets, int *csr_colindices, double *csr_values,
142                        double const *x, double *y)
143 {
144     for (size_t i=0; i<N; ++i) {
145         double value = 0;
146         for (size_t j=csr_rowoffsets[i]; j<csr_rowoffsets[i+1]; ++j)
147             value += csr_values[j] * x[csr_colindices[j]];
148         y[i] = value;
149     }
150 }
151
152
153 double benchmark_ocl(size_t N, size_t max_nonzeros_per_row,
154                     int* csr_rowoffsets, int *csr_colindices,
155                     double* csr_values,
156                     double* x, double* y)
157 {
158     cl_int err;
159     Timer timer1;
160     //
161     ////////////////////////////////////////////////// Part 1: Set up an OpenCL context with one device
162     //////////////////////////////////////////////////
163     //
164     //
165     // Query platform:
166     //
167     cl_uint num_platforms;
168     cl_platform_id platform_ids[42]; //no more than 42 platforms supported...
169     err = clGetPlatformIDs(42, platform_ids, &num_platforms); OPENCL_ERR_CHECK(err);
170     //std::cout << "# Platforms found: " << num_platforms << std::endl;
171     //
172     //
173     // Query devices:
174     //
175     cl_device_id device_ids[42];

```

```

176     cl_uint num_devices;
177     char device_name[64];
178     cl_device_id my_device_id;
179     cl_platform_id my_platform;
180     for (int i = 0; i < num_platforms; ++i)
181     {
182         my_platform = platform_ids[i];
183         if (target == "GPU") {
184             err = clGetDeviceIDs(my_platform, CL_DEVICE_TYPE_GPU, 42, device_ids, &num_devices);
185         }
186         else {
187             err = clGetDeviceIDs(my_platform, CL_DEVICE_TYPE_CPU, 42, device_ids, &num_devices);
188         }
189         if (err == CL_SUCCESS)
190             break;
191     }
192     OPENCL_ERR_CHECK(err);
193     //std::cout << "# Devices found: " << num_devices << std::endl;
194     my_device_id = device_ids[0];
195
196     size_t device_name_len = 0;
197     err = clGetDeviceInfo(my_device_id, CL_DEVICE_NAME, sizeof(char)*63, device_name, &
        device_name_len); OPENCL_ERR_CHECK(err);
198
199     std::cout << "Using the following device: " << device_name << std::endl;
200
201     //
202     // Create context:
203     //
204     cl_context my_context = clCreateContext(0, 1, &my_device_id, NULL, NULL, &err);
        OPENCL_ERR_CHECK(err);
205
206     //
207     // create a command queue for the device:
208     //
209     cl_command_queue my_queue = clCreateCommandQueueWithProperties(my_context, my_device_id,
        0, &err); OPENCL_ERR_CHECK(err);
210
211     //
212     ////////////////////////////////////////////////// Part 2: Create a program and extract kernels
        //////////////////////////////////////////
213     //
214     //
215     // Build the program:
216     //
217     const char* ocl_prog = my_opencl_program.c_str();
218     size_t source_len = my_opencl_program.length();
219
220     cl_program prog = clCreateProgramWithSource(my_context, 1, &ocl_prog, &source_len, &err);
        OPENCL_ERR_CHECK(err);
221     err = clBuildProgram(prog, 0, NULL, NULL, NULL, NULL);
222
223     //
224     // Print compiler errors if there was a problem:
225     //
226     if (err != CL_SUCCESS) {
227
228         char *build_log;
229         size_t ret_val_size;
230         err = clGetProgramBuildInfo(prog, my_device_id, CL_PROGRAM_BUILD_LOG, 0, NULL, &
            ret_val_size);
231         build_log = (char *)malloc(sizeof(char) * (ret_val_size+1));
232         err = clGetProgramBuildInfo(prog, my_device_id, CL_PROGRAM_BUILD_LOG, ret_val_size,
            build_log, NULL);
233         build_log[ret_val_size] = '\0'; // terminate string
234         std::cout << "Log: " << build_log << std::endl;
235         free(build_log);
236         std::cout << "OpenCL program sources: " << std::endl << my_opencl_program << std::endl;
237         return EXIT_FAILURE;
238     }
239     //
240     // Extract the only kernel in the program:

```

```

241 //
242
243 #ifdef SLOW_KERNEL
244     cl_kernel my_kernel = clCreateKernel(prog, "ocl_csr_matvec", &err); OPENCL_ERR_CHECK(err);
245 #endif
246 #ifdef FAST_KERNEL
247     cl_kernel my_kernel = clCreateKernel(prog, "ocl_csr_matvec", &err); OPENCL_ERR_CHECK(err);
248 #endif
249
250 //
251 ////////////////////////////////////////////////// Part 3: Create memory buffers
252 //////////////////////////////////////////////////
253 //
254 cl_uint vector_size = N;
255 size_t local_size = LOCAL_SIZE;
256 size_t global_size = GLOBAL_SIZE*GLOBAL_SIZE;
257 size_t groups = 1 + int(N/LOCAL_SIZE);
258
259 cl_mem ocl_x = clCreateBuffer(my_context, CL_MEM_READ_WRITE | CL_MEM_COPY_HOST_PTR, sizeof
    (ScalarType)*N, x, &err); OPENCL_ERR_CHECK(err);
260
261 cl_mem ocl_y = clCreateBuffer(my_context, CL_MEM_READ_WRITE | CL_MEM_COPY_HOST_PTR, sizeof
    (ScalarType)*N, y, &err); OPENCL_ERR_CHECK(err);
262
263 cl_mem ocl_csr_rowoffsets = clCreateBuffer(my_context, CL_MEM_READ_WRITE |
    CL_MEM_COPY_HOST_PTR, sizeof(int)*(N+1), csr_rowoffsets, &err); OPENCL_ERR_CHECK(err);
264
265 cl_mem ocl_csr_colindices = clCreateBuffer(my_context, CL_MEM_READ_WRITE |
    CL_MEM_COPY_HOST_PTR, sizeof(int)*(max_nonzeros_per_row*N), csr_colindices, &err);
    OPENCL_ERR_CHECK(err);
266
267 cl_mem ocl_csr_values = clCreateBuffer(my_context, CL_MEM_READ_WRITE |
    CL_MEM_COPY_HOST_PTR, sizeof(double)*(max_nonzeros_per_row*N), csr_values, &err);
    OPENCL_ERR_CHECK(err);
268
269 //
270 ////////////////////////////////////////////////// Part 4: Run kernel //////////////////////////////////////
271 //
272 //
273 // Set kernel arguments:
274 //
275
276 // "__kernel void ocl_csr_matvec(size_t N,\n"
277 // "                __global int *csr_rowoffsets,\n"
278 // "                __global int *csr_colindices,\n"
279 // "                __global double *csr_values,\n"
280 // "                __global double const *x, __global double *y)\n"
281 err = clSetKernelArg(my_kernel, 0, sizeof(cl_uint), (void*)&vector_size); OPENCL_ERR_CHECK
    (err);
282 err = clSetKernelArg(my_kernel, 1, sizeof(cl_mem), (double*)&ocl_csr_rowoffsets);
    OPENCL_ERR_CHECK(err);
283 err = clSetKernelArg(my_kernel, 2, sizeof(cl_mem), (double*)&ocl_csr_colindices);
    OPENCL_ERR_CHECK(err);
284 err = clSetKernelArg(my_kernel, 3, sizeof(cl_mem), (double*)&ocl_csr_values);
    OPENCL_ERR_CHECK(err);
285 err = clSetKernelArg(my_kernel, 4, sizeof(cl_mem), (double*)&ocl_x); OPENCL_ERR_CHECK(err
    );
286 err = clSetKernelArg(my_kernel, 5, sizeof(cl_mem), (double*)&ocl_y); OPENCL_ERR_CHECK(err
    );
287
288 //
289 // Enqueue kernel in command queue:
290 //
291 timer1.reset();
292 err = clEnqueueNDRangeKernel(my_queue, my_kernel, 1, NULL, &global_size, &local_size, 0,
    NULL, NULL); OPENCL_ERR_CHECK(err);
293
294 // wait for all operations in queue to finish:
295 err = clFinish(my_queue); OPENCL_ERR_CHECK(err);
296 double ocl_time = timer1.get();
297

```

```

298 //
299 ////////////////////////////////////////////////// Part 5: Get data from OpenCL buffer
    //////////////////////////////////////
300 //
301 err = clEnqueueReadBuffer(my_queue, ocl_y, CL_TRUE, 0, sizeof(ScalarType)*vector_size, y,
    0, NULL, NULL); OPENCL_ERR_CHECK(err);
302
303 // wait for all operations in queue to finish:
304 err = clFinish(my_queue); OPENCL_ERR_CHECK(err);
305 //
306 // cleanup
307 //
308 clReleaseMemObject(ocl_x);
309 clReleaseMemObject(ocl_y);
310 clReleaseMemObject(ocl_csr_rowoffsets);
311 clReleaseMemObject(ocl_csr_colindices);
312 clReleaseMemObject(ocl_csr_values);
313 clReleaseProgram(prog);
314 clReleaseCommandQueue(my_queue);
315 clReleaseContext(my_context);
316
317 std::cout << "From OCL benchmark: " << ocl_time << std::endl;
318
319 return ocl_time;
320 }
321
322
323 /** Solve a system with 'points_per_direction * points_per_direction' unknowns */
324 void benchmark_matvec(size_t points_per_direction, size_t max_nonzeros_per_row,
325     void (*generate_matrix)(size_t, int*, int*, double*),
326     std::string gen_type) // function pointer parameter
327 {
328
329     size_t N = points_per_direction * points_per_direction; // number of rows and columns
330     std::fstream csv;
331     csv.open(csv_name, std::fstream::out | std::fstream::app);
332
333     //
334     // Allocate CSR arrays.
335     //
336     // Note: Usually one does not know the number of nonzeros in the system matrix a-priori.
337     // For this exercise, however, we know that there are at most 5 nonzeros per row in
    the system matrix, so we can allocate accordingly.
338     //
339     int *csr_rowoffsets = (int*)malloc(sizeof(double) * (N+1));
340     int *csr_colindices = (int*)malloc(sizeof(double) * max_nonzeros_per_row * N);
341     double *csr_values = (double*)malloc(sizeof(double) * max_nonzeros_per_row * N);
342
343     //
344     // fill CSR matrix with values
345     //
346     generate_matrix(points_per_direction, csr_rowoffsets, csr_colindices, csr_values);
347
348     //
349     // Allocate vectors:
350     //
351     double *x = (double*)malloc(sizeof(double) * N); std::fill(x, x + N, 1);
352     double *y = (double*)malloc(sizeof(double) * N); std::fill(y, y + N, 0);
353     double *y_ocl = (double*)malloc(sizeof(double) * N); std::fill(y, y + N, 0);
354
355     //
356     // Call matrix-vector product reference
357     //
358     Timer timer;
359     timer.reset();
360     csr_matvec_product(N, csr_rowoffsets, csr_colindices, csr_values, x, y);
361     double cpu_time = timer.get();
362
363     //
364     // Call matrix-vector product kernel
365     //
366     double ocl_time = benchmark_ocl(N, max_nonzeros_per_row, csr_rowoffsets, csr_colindices,

```



```

        csr_values, x, y_ocl);
367     std::cout << "Reference: " << std::endl;
368     printContainer(y, N, 10);
369     std::cout << "OpenCL: " << std::endl;
370     printContainer(y_ocl, N, 10);
371     double difference = diff_norm(y_ocl, y, N);
372     bool check_res = check(y_ocl, y, N);
373     std::cout << "Difference between the two: " << difference << " (check: " << check_res <<
        std::endl;
374     std::cout << "Time for ref product: " << cpu_time << std::endl;
375     std::cout << "Time for OCL product: " << ocl_time << std::endl;
376
377     csv << N << ";";
378     << points_per_direction << ";";
379     << target << ";";
380     << gen_type << ";";
381     << ocl_time << ";";
382     << cpu_time << ";";
383     << difference << ";";
384     << check_res << std::endl;
385
386     csv.close();
387     free(x);
388     free(y);
389     free(y_ocl);
390     free(csr_rowoffsets);
391     free(csr_colindices);
392     free(csr_values);
393 }
394
395
396 int main() {
397     std::fstream csv;
398     #ifdef TRUNC_CSV
399     csv.open(csv_name, std::fstream::out | std::fstream::trunc);
400     csv << "N;points_per_direction;target;gen_type;ocl_time;cpu_time;check;diff_norm" << std::
        endl;
401     #endif
402     #ifndef TRUNC_CSV
403     csv.open(csv_name, std::fstream::out | std::fstream::app);
404     #endif
405
406     uint pp = PP;
407     std::cout << "# Benchmarking finite difference matrix" << std::endl;
408     benchmark_matvec(pp, 5, generate_fdm_laplace, "1"); // 100*100 unknowns, finite difference
        matrix
409     // the last string is just so that I know which matrix was used
410
411     std::cout << "# Benchmarking special matrix" << std::endl;
412     benchmark_matvec(pp, 2000, generate_matrix2, "2"); // 100*100 unknowns, special matrix
        with 200-2000 nonzeros per row
413
414     std::cout << "Data: \nRuntimes in csv form can be found here\nhttps://gtx1080.360252.org
        /2020/ex7/" << csv_name << std::endl;
415
416     return EXIT_SUCCESS;
417 }

```